

Relational Model and Relational Algebra

CMPSCI 445 – Database Systems
Fall 2008

Next lectures: Querying relational data

- Today
 - Relational model, Relational algebra
- Wednesday
 - SQL
- Homework 1 assigned Friday
- Next Week
 - More SQL

The Relational Model

- The **relational data model** (Codd, 1970):
 - **Data independence**: details of physical storage are hidden from users
 - High-level **declarative query language**
 - say **what** you want, not **how** to compute it.
 - mathematical foundation

Relational Database: Definitions

Relational Database: Definitions

- *Relational database*: a set of *relations*

Relational Database: Definitions

- *Relational database*: a set of *relations*
- *Relation*: made up of 2 parts:
 - *Schema* : specifies name of relation, plus name and type/domain of each column.
 - *Instance* : a *table*, with rows and columns.

Relational Database: Definitions

- *Relational database*: a set of *relations*
- *Relation*: made up of 2 parts:
 - *Schema* : specifies name of relation, plus name and type/domain of each column.
 - *Instance* : a *table*, with rows and columns.

Students(*sid*: string, *name*: string, *login*: string, *age*: integer, *gpa*: real).

Restriction: all attributes are of atomic type,
no nested tables

Relational instances: tables

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Relational instances: tables

Students

column,
attribute,
field

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

Relational instances: tables

Students

column,
attribute,
field

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

row, tuple

Relational instances: tables

Students

column,
attribute,
field

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

row, tuple

Attribute value

Relational instances: tables

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@eecs	18	3.2
53650	Smith	smith@math	19	3.8

column, attribute, field

row, tuple

Attribute value

A relation is a **set** of tuples: no tuple can occur more than once

- Real systems may allow duplicates for efficiency or other reasons – we'll come back to this.

Example Database

STUDENT

sid	name
1	Jill
2	Bo
3	Maya

Takes

sid	cid
1	445
1	483
3	435

COURSE

cid	title	sem
445	DB	F08
483	AI	S08
435	Arch	F08

PROFESSOR

fid	name
1	Diao
2	Saul
8	Weems

Teaches

fid	cid
1	445
2	483
8	435

Relational Query Languages

- Query languages: Allow manipulation and retrieval of data from a database.
- DB query languages **!=** programming languages
 - not expected to be “Turing complete”.
 - not intended to be used for complex calculations.
 - support easy, efficient access to large data sets.

Query language preliminaries

Query $Q: R_1..R_n \rightarrow R'$

- A query is applied to one or more *relation instances*
- The result of a query is a relation instance.
- Input and output schema:
 - *Schema of input* relations for a query are *fixed*
 - The *schema for the result* of a given query is also *fixed*: determined by definition of query language constructs.

What is an “Algebra” ?

What is an “Algebra” ?

- Mathematical system consisting of:

What is an “Algebra” ?

- Mathematical system consisting of:
 - *Operands* --- variables or values from which new values can be constructed.

What is an “Algebra” ?

- Mathematical system consisting of:
 - *Operands* --- variables or values from which new values can be constructed.
 - *Operators* --- symbols denoting procedures that construct new values from given values.

What is the Relational Algebra?

- An algebra whose operands are relations or variables that represent relations.
- Operators are designed to do the most common things that we need to do with relations in a database.
 - The result is an algebra that can be used as a *query language* for relations.

Relational Algebra

- Operates on relations, i.e. *sets*
 - Later: we discuss how to extend this to *bags*
- Five basic operators:
 - Union: \cup
 - Difference: $-$
 - Selection: σ
 - Projection: Π
 - Cartesian Product: \times
- Derived or auxiliary operators:
 - Intersection, complement
 - Joins (natural, equi-join, theta join)
 - Renaming: ρ

1. Union and 2. Difference

R_1

sid	name
1	Jill
2	Bo
3	Maya

R_2

sid	name
1	Jill
4	Bob

$R_1 \cup R_2$



sid	name
1	Jill
2	Bo
3	Maya
4	Bob

$R_1 - R_2$

sid	name
2	Bo
3	Maya

What about Intersection ?

- It is a derived operator
- $R_1 \cap R_2 = R_1 - (R_1 - R_2)$
- Also expressed as a join (will see later)



3. Selection

- Returns all tuples which satisfy a condition

- Notation: $\sigma_c(R)$

- Examples

$\sigma_{\text{CID} > 400}(\text{Course})$

$\sigma_{\text{title} = \text{"AI"}}(\text{Course})$

- The condition c can be =, <, ≤, >, ≥, <>

Course

cid	title	sem
445	DB	F08
483	AI	S08
435	Arch	F08

4. Projection

- Eliminates columns, then removes duplicates
- Notation: $\Pi_{A_1, \dots, A_n}(R)$
- Example: project cid and name

$\Pi_{\text{cid, name}}(\mathbf{Course})$

Output schema: **Answer(cid, name)**

4. Projection

- Eliminates columns, then removes duplicates
- Notation: $\Pi_{A_1, \dots, A_n}(R)$
- Example: project cid and name

$\Pi_{\text{cid, name}}(\mathbf{Course})$

Output schema: **Answer(cid, name)**

Course

cid	name	sem
445	DB	F08
483	AI	S08
445	DB	S08



Answer

cid	name
445	DB
483	AI

5. Cartesian Product

- Each tuple in R_1 with each tuple in R_2
- Notation: $R_1 \times R_2$
- Very rare in practice; mainly used to express joins

Also called “Cross Product”

Cartesian Product

Student

sid	name
1	Jill
2	Bo

Takes

sid	cid
1	445
1	483
3	435

Cartesian Product

Student

sid	name
1	Jill
2	Bo

Takes

sid	cid
1	445
1	483
3	435

Student \times Takes

sid	name	sid	cid
1	Jill	1	445
1	Jill	1	483
1	Jill	3	435
2	Bo	1	445
2	Bo	1	483
2	Bo	3	435

Renaming

- Changes the **schema**, not the **instance**
- Notation: $\rho_{B1, \dots, Bn}(R)$
- Example:
 $\rho_{\text{courseID}, \text{cname}, \text{term}}(\mathbf{Course})$

Renaming

- Changes the **schema**, not the **instance**
- Notation: $\rho_{B1, \dots, Bn} (R)$
- Example:

$\rho_{\text{courseID}, \text{cname}, \text{term}} (\text{Course})$

Course

cid	name	sem
445	DB	F08
483	AI	S08
445	DB	S08



courseID	cname	term
445	DB	F08
483	AI	S08
445	DB	S08

Natural Join

- Notation: $R_1 \bowtie R_2$
- Meaning: $R_1 \bowtie R_2 = \Pi_A(\sigma_C(R_1 \times R_2))$
- Where:
 - The selection σ_C checks equality of all common attributes
 - The projection eliminates the **duplicate** common attributes

Natural join example

Student

sid	name
1	Jill
2	Bo
3	Maya

Calculate:

Student \bowtie Takes

Takes

sid	cid
1	445
1	483
3	435

Natural join example

Student

sid	name
1	Jill
2	Bo
3	Maya


Takes

sid	cid
1	445
1	483
3	435

Calculate:

Student \bowtie Takes

sid	name	cid
1	Jill	445
1	Jill	483
3	Maya	435



Theta Join

- A join that involves a predicate
- $R1 \bowtie_{\theta} R2 = \sigma_{\theta} (R1 \times R2)$
- Here θ can be any condition:
 $=, <, \neq, \leq, >, \geq$

Example: Student $\bowtie_{\text{age} > \text{age}}$ Prof

Equi-join


- A theta join where θ is an equality
- $R_1 \bowtie_{A=B} R_2 = \sigma_{A=B} (R_1 \times R_2)$
- Very useful join in practice
- Example: Student $\bowtie_{\text{sid}=\text{sid}}$ Takes

Review

- Five basic operators of the Relational Algebra:
 - Union: \cup
 - Difference: $-$
 - Selection: σ
 - Projection: Π
 - Cartesian Product: \times
- Derived or auxiliary operators:
 - Intersection, complement
 - Joins (natural, equi-join, theta join)
 - Renaming: ρ

Combining operators: complex expressions

$\Pi_{\text{name,sid}} (\sigma_{\text{name}=\text{"DB"}} (\text{Course} \bowtie (\text{Students} \bowtie \text{Takes})))$



In-class exercise

- Please calculate:

$\Pi_{\text{name, sid}} (\sigma_{\text{title}=\text{"DB"}} (\text{Course} \bowtie (\text{Students} \bowtie \text{Takes})))$

Course

cid	title	sem
445	DB	F08
483	AI	S08
435	Arch	F08

Students

sid	name
1	Jill
2	Bo
3	Maya

Takes

sid	cid
1	445
1	483
3	435

Query equivalence

Definition: **Query Equivalence**

Two queries Q and Q' are equivalent if:

for all databases D , $Q(D) = Q'(D)$

Query equivalence

Query Optimization

Is Based on Algebraic Equivalences

- Relational algebra has laws of commutativity, associativity, etc. that imply certain expressions are **equivalent**.
- They may be different in cost of evaluation!

$$\sigma_{c \wedge d}(R) \equiv \sigma_c(\sigma_d(R)) \quad \text{cascading selection}$$

$$R \bowtie (S \bowtie T) \equiv (R \bowtie S) \bowtie T \quad \text{join associativity}$$

$$\sigma_c(R \bowtie S) \equiv \sigma_c(R) \bowtie S \quad \text{pushing selections}$$

- Query optimization finds the most efficient representation to evaluate (or one that's not bad)

Relational calculus

- What is a “calculus”?
 - The term "calculus" means a system of computation
 - The relational calculus is a system of computing with relations

Relational calculus (in 1 slide)

English: Name and sid of students who are taking the
 course “DB”

Relational calculus (in 1 slide)

English: Name and sid of students who are taking the course "DB"

RA: $\Pi_{\text{name,sid}} (\text{Students} \bowtie \text{Takes} \bowtie \sigma_{\text{name}=\text{"DB"}} (\text{Course}))$

Relational calculus (in 1 slide)

English: Name and sid of students who are taking the course "DB"

RA: $\Pi_{\text{name,sid}} (\text{Students} \bowtie \text{Takes} \bowtie \sigma_{\text{name}=\text{"DB"}} (\text{Course}))$

RC: $\{x_{\text{name}}, x_{\text{sid}} \mid \exists x_{\text{cid}} \exists x_{\text{term}} \text{Students}(x_{\text{sid}}, x_{\text{name}}) \wedge \text{Takes}(x_{\text{sid}}, x_{\text{cid}}) \wedge \text{Course}(x_{\text{cid}}, \text{"DB"}, x_{\text{term}}) \}$

Relational calculus (in 1 slide)

English: Name and sid of students who are taking the course "DB"

RA: $\Pi_{\text{name,sid}} (\text{Students} \bowtie \text{Takes} \bowtie \sigma_{\text{name}=\text{"DB"}} (\text{Course}))$

RC: $\{ \underline{x_{\text{name}}}, \underline{x_{\text{sid}}} \mid \exists \underline{x_{\text{cid}}} \exists \underline{x_{\text{term}}} \text{Students}(\underline{x_{\text{sid}}}, \underline{x_{\text{name}}}) \wedge \text{Takes}(\underline{x_{\text{sid}}}, \underline{x_{\text{cid}}}) \wedge \text{Course}(\underline{x_{\text{cid}}}, \text{"DB"}, \underline{x_{\text{term}}}) \}$

Where are the joins?

Algebra v. Calculus

- Relational Algebra: More operational; very useful for representing execution plans.
- Relational Calculus: More declarative, basis of SQL
- The calculus and algebra have equivalent expressive power (Codd)

A language that can express this core class of queries is called **Relationally Complete**