

# SQL Overview

- Query capabilities
  - SELECT-FROM-WHERE blocks,
  - Basic features, ordering, duplicates
  - Set operations (union, intersect, except)
  - Aggregation & Grouping
  - Nested queries (correlation)
  - Null values

# Set operations

- UNION
- INTERSECTION
- EXCEPT (sometimes called MINUS)
- Recall: schemas must match for these operations.

# UNION example

*Find the names of sailors who have reserved a red or a green boat.*

```
SELECT sname  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
```

UNION

```
SELECT sname  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'green'
```

# UNION

- Duplicates **ARE NOT** eliminated by default in basic SELECT-FROM-WHERE queries
- Duplicate **ARE** eliminated by default for UNION queries.
- To preserve duplicates in UNION, you must use UNION ALL

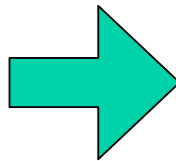
## UNION example, alternative:

*Find the names of sailors who have reserved a red or a green boat.*

```
SELECT DISTINCT sname  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid  
      AND (B.color = 'red' OR B.color = 'green')
```

## A small change in this query...

*Find the names of sailors  
who have reserved a red  
**or** a green boat.*



*Find the names of sailors  
who have reserved a red  
**and** a green boat.*

```
SELECT DISTINCT sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
      AND (B.color = 'red' OR B.color = 'green')
```

```
SELECT sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid
      AND (B.color = 'red' AND B.color = 'green')
```

**This doesn't work! What  
does this query return?**

*Find the names of sailors who have reserved a red  
**and** a green boat.*

```
SELECT sname  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
```

INTERSECT

```
SELECT sname  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'green'
```

# SQL Overview

- Query capabilities
  - SELECT-FROM-WHERE blocks,
  - Basic features, ordering, duplicates
  - Set ops (union, intersect, except)
  - Aggregation & Grouping
  - Nested queries (correlation)
  - Null values



# Aggregation

```
SELECT Avg(S.age)
FROM Sailors
WHERE S.rating = 10
```

SQL supports several aggregation operations:

```
COUNT (*)
COUNT ( [DISTINCT] A)
SUM ( [DISTINCT] A)
AVG ( [DISTINCT] A)
MAX (A)
MIN (A)
```

# Aggregation: Count

```
SELECT Count(*)  
FROM Sailors  
WHERE rating > 5
```

Except for COUNT, all aggregations apply to a single attribute

# Aggregation: Count

COUNT applies to duplicates, unless otherwise stated:

```
SELECT Count(category)
FROM Product
WHERE year > 1995
```

Better:

```
SELECT Count(DISTINCT category)
FROM Product
WHERE year > 1995
```

# Simple Aggregation

Purchase(product, date, price, quantity)

Example 1: find total sales for the entire database

```
SELECT Sum(price * quantity)
FROM Purchase
```

Example 1': find total sales of bagels

```
SELECT Sum(price * quantity)
FROM Purchase
WHERE product = 'bagel'
```

# GROUP BY and HAVING clauses

- We often want to apply aggregates to each of a number of groups of rows in a relation.

*Find the age of the youngest sailor for each rating level.*

```
SELECT MIN (S.age)
FROM   Sailors S
WHERE  S.rating = i
```

For  $i = 1, 2, \dots, 10$

# Grouping

Sailors

sid	sname	rating	age
29	brutus	1	33
85	art	3	25.5
95	bob	3	63.5
96	frodo	3	25.5
22	dustin	7	45
64	horatio	7	35
31	lubber	8	55.5
32	andy	8	25.5
74	horatio	9	35
58	rusty	10	35
71	zorba	10	16

```
SELECT S.rating, MIN(S.age)
FROM Sailors S
GROUP BY S.rating
```

New Table

rating	age?
1	
3	
7	
8	
9	
10	

# Queries With GROUP BY and HAVING

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>

- The *target-list* contains (i) attribute names (ii) **terms with aggregate operations** (e.g., MIN (*S.age*)).
  - The attribute list (i) must be a subset of *grouping-list*. Intuitively, each answer tuple corresponds to a *group*, and these attributes must have a single value per group.

# Conceptual Evaluation

- The cross-product of ***relation-list*** is computed, tuples that fail ***qualification*** are discarded, `unnecessary' fields are deleted, and the remaining tuples are partitioned into groups by the value of attributes in ***grouping-list***.
- The ***group-qualification*** is then applied to eliminate some groups. Expressions in *group-qualification* must have a ***single value per group!***
- One answer tuple is generated per qualifying group.



*Find age of the youngest sailor with age  $\geq 18$ , for each rating with at least 2 such sailors*

```
SELECT S.rating, MIN (S.age)
      AS minage
FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING COUNT (*) > 1
```

*Answer relation:*

rating	minage
3	25.5
7	35.0
8	25.5

*Sailors instance:*

sid	sname	rating	age
22	dustin	7	45.0
29	brutus	1	33.0
31	lubber	8	55.5
32	andy	8	25.5
58	rusty	10	35.0
64	horatio	7	35.0
71	zorba	10	16.0
74	horatio	9	35.0
85	art	3	25.5
95	bob	3	63.5
96	frodo	3	25.5

65

# SQL Overview

- Query capabilities
  - SELECT-FROM-WHERE blocks,
  - Basic features, ordering, duplicates
  - Set ops (union, intersect, except)
  - Aggregation & Grouping
  - Nested queries (correlation)
  - Null values

# Nested queries

- A **nested query** is a query with another query embedded within it.
- The embedded query is called the **subquery**.
- The subquery usually appears in the WHERE clause:

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.sid IN ( SELECT R.sid
                   FROM Reserves R
                   WHERE R.bid = 103 )
```

(Subqueries also possible in FROM or HAVING clause.)

# Conceptual evaluation, extended

- For each row in cross product of outer query, evaluate the WHERE clause conditions, (re)computing the subquery.

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.sid IN ( SELECT R.sid
                  FROM Reserves R
                  WHERE R.bid = 103 )
```

equivalent to:

```
SELECT  S.sname
FROM    Sailors S, Reserves R
WHERE   S.sid=R.sid AND R.bid=103
```

# Correlated subquery

- If the inner subquery depends on tables mentioned in the outer query then it is a **correlated subquery**.
- In terms of conceptual evaluation, we must recompute subquery for each row of outer query.

```
SELECT  S.sname
FROM    Sailors S
WHERE   EXISTS ( SELECT *
                  FROM    Reserves R
                  WHERE   R.bid = 103
                  AND R.sid = S.sid )
```

Correlation



# Set-comparison operators

- Optional NOT may precede these:
  - **EXISTS R** -- true if R is non-empty
  - **attr IN R** -- true if R contains attr
  - **UNIQUE R** -- true if no duplicates in R
- For arithmetic operator **op** {<, <=, =, < >, >=, >}
  - **attr op ALL R** -- all elements of R satisfy condition
  - **attr op ANY R** -- some element of R satisfies condition

**IN** equivalent to **= ANY**

**NOT IN** equivalent to **< > ALL**

# Example

- Find the sailors with the highest rating

```
SELECT S.sid  
FROM   Sailors S  
WHERE  S.rating >= ALL (SELECT S2.rating  
                        FROM Sailors S2 )
```

# Please write SQL

- Find sailors whose rating is higher than **some** sailor named Horatio.

```
SELECT S.sid
FROM   Sailors S
WHERE  S.rating > ANY (SELECT S2.rating
                       FROM Sailors S2
                       WHERE S2.name = 'Horatio')
```

- Find sailors whose rating is higher than **all** sailors named Horatio.

```
SELECT S.sid
FROM   Sailors S
WHERE  S.rating > ALL (SELECT S2.rating
                      FROM Sailors S2
                      WHERE S2.name = 'Horatio')
```



Find boats **not** reserved by sailor with  
sid = 100.

- B: all boats
- R: boats reserved by sailor with sid=100
- B – R is what we want.

```
SELECT B.bid  
FROM Boats B  
WHERE B.bid NOT IN (SELECT R.bid  
                     FROM Reserves R  
                     WHERE R.sid = 100 );
```

# Existential conditions

- Find the names of sailors who have reserved **some** boat
- (i.e. *there exists* a boat they reserved)

```
SELECT S.sname  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid
```

- Existential conditions are natural and easy.

# Universal conditions

- Find the names of sailors who have reserved **all** boats.
- (i.e. *for each* boat, they have reserved it.)
- Universal conditions are harder.

```
SELECT S.sname
FROM   Sailors S
WHERE  NOT EXISTS (
                Set of boats not
                reserved by S.sid
            )
```

# Universal conditions

- Find the names of sailors who have reserved **all** boats.

```
SELECT S.sname  
FROM   Sailors S  
WHERE  NOT EXISTS (  
  
      )
```

boats the sailor hasn't reserved  
(we just wrote this query)



```
SELECT B.bid  
FROM   Boats B  
WHERE  B.bid NOT IN (SELECT R.bid  
                    FROM   Reserves R  
                    WHERE  R.sid = S.sid )
```

For each sailor, check that there is no boat s/he hasn't reserved.

# Simulating INTERSECT

- Suppose we have tables R(a,b) and S(a,b)
- The following computes  $R \cap S$ :

```
SELECT DISTINCT *  
FROM R  
WHERE (R.a, R.b) IN (SELECT *  
                     FROM S );
```

This can be expressed without nesting:

- Given R(a,b), S(a,b),  
what is  $R \bowtie S$  ?

Intersection!

```
SELECT DISTINCT R.a, R.b  
FROM R, S  
WHERE R.a = S.a AND R.b = S.b;
```

*Find the names of sailors who reserved a **red** and a **green** boat.*

using INTERSECT

```
SELECT sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
```

**INTERSECT**

```
SELECT sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'green'
```

without INTERSECT

```
SELECT sname
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = 'red'
      AND S.sid IN
      (SELECT S2.sid
       FROM Sailors S2, Reserves R2, Boats B2
       WHERE S2.sid = R2.sid AND R2.bid = B2.bid AND B2.color = 'green' )
```

“Find all sailors who have reserved a red boat and, further, have **sids** that are included in the set of **sids** of sailors who have reserved a green boat.”

# Simulating EXCEPT (set difference)

- What does this query compute?

```
SELECT DISTINCT *  
FROM R  
WHERE (R.a, R.b) NOT IN (SELECT *  
                        FROM S );
```

Can this be expressed without a nested query? No.

(But this fact is not obvious)