

Storage

CMPSCI 445

Fall 2008

Disks and DBMS Design

- ❖ DBMS stores information on disks.
- ❖ This has major implications for DBMS design!
 - **READ:** transfer data from disk to main memory (RAM) for data processing.
 - **WRITE:** transfer data from RAM to disk for persistent storage.
 - Both are high-cost operations, relative to in-memory operations, so must be planned carefully!

Why Not Store Everything in Main Memory?


- ❖ *Main memory is volatile.* We want data to be saved between runs. (Obviously!)
- ❖ *Costs too much.* \$100 will buy you either 1GB of RAM or 160GB of disk today.
- ❖ *32-bit addressing limitation.*
 - 2^{32} bytes can be directly addressed in memory.
 - Number of objects cannot exceed this number.

Basics of Disks

- ❖ Unit of storage and retrieval: *disk block* or *page*.
 - A disk block/page is a contiguous sequence of bytes.
 - Size of a DBMS parameter, 4KB or 8KB.
- ❖ Disks support direct access to a page.
- ❖ Unlike RAM, time to retrieve a page varies!
 - It depends upon the location on disk.
 - Therefore, relative placement of pages on disk has major impact on DBMS performance!

Components of a Disk

- ❖ Platters spin (say, 7200rpm).
- ❖ Arm assembly is moved in or out to position a head on a desired *track*.
- ❖ Only one head reads/writes at any one time.
- ❖ Tracks under heads make a *cylinder* (imaginary!).
- ❖ Each track is divided into *sectors* (whose size is fixed).
- ❖ *Block size* is a multiple of *sector size*.



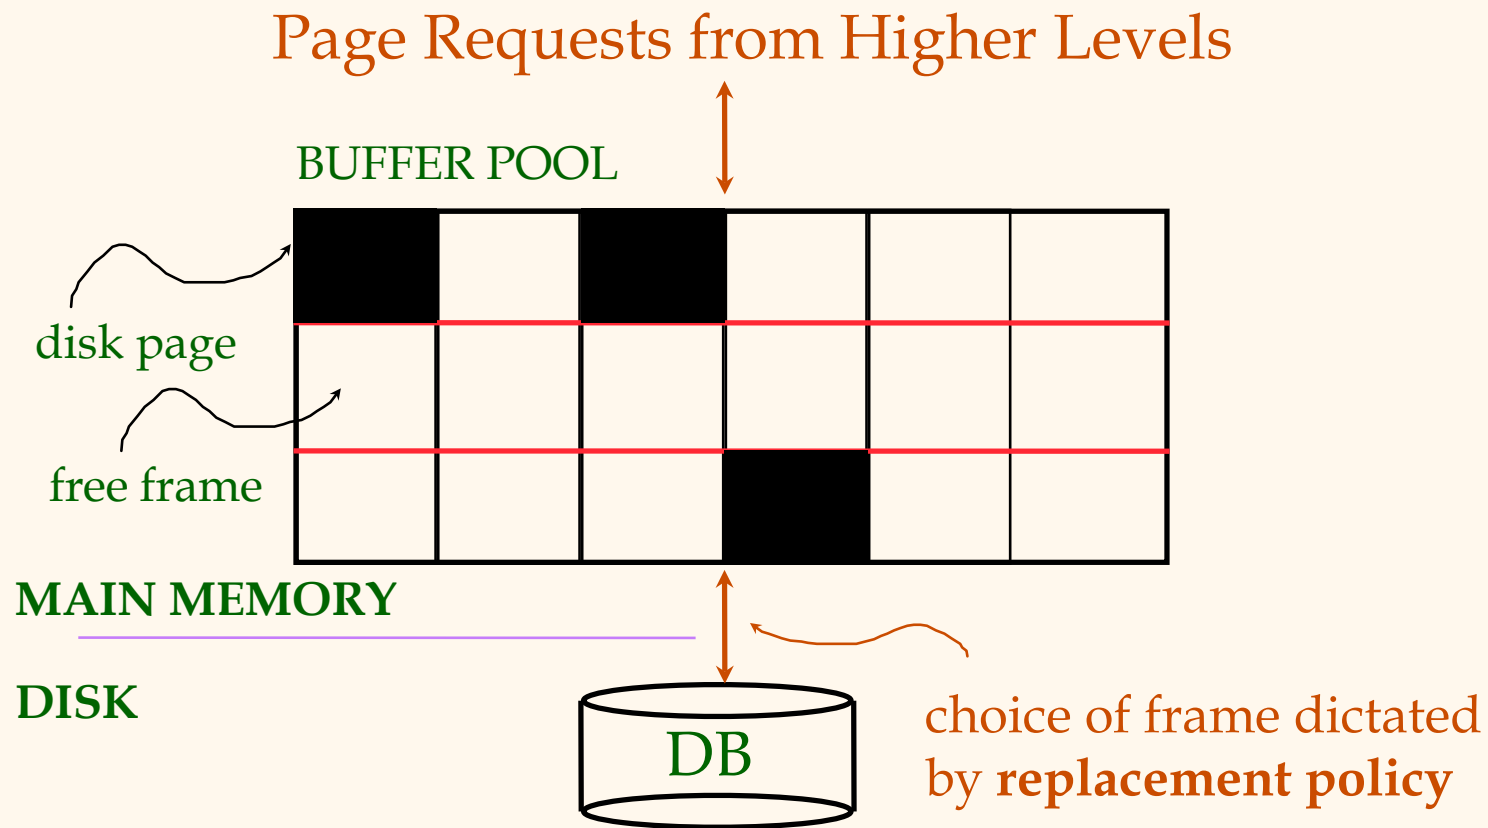
Accessing a Disk Page

- ❖ Time to access (read/write) a disk block:
 - *seek time* (moving arms to position disk head on track)
 - *rotational delay* (waiting for block to rotate under head)
 - *transfer time* (actually moving data to/from disk surface)
- ❖ Seek time and rotational delay dominate.
 - Seek time varies from about 1 to 20msec
 - Rotational delay varies from 0 to 10msec
 - Transfer rate is about 1msec per 4KB page
- ❖ Key to lower I/O cost: **reduce seek/rotation delays!**

Arranging Pages on Disk

- ❖ *'Next'* block concept:
 - blocks on same track, followed by
 - blocks on same cylinder, followed by
 - blocks on adjacent cylinder
- ❖ Blocks in a file should be arranged sequentially on disk (by *'next'*), to minimize seek and rotational delay.
- ❖ For a *sequential scan*, *pre-fetching* several pages at a time is a big win!

Buffer Management in a DBMS



- ❖ *Data must be in RAM for DBMS to operate on it!*
- ❖ *Table of $\langle \text{frame\#}, \text{pageid} \rangle$ pairs is maintained.*

More on Buffer Management

- ❖ Requestor of page must unpin it, and indicate whether page has been modified:
 - *dirty* bit is used for this.
- ❖ Page in pool may be requested many times,
 - a *pin count* is used. A page is a candidate for replacement iff *pin count* = 0.
- ❖ CC & recovery may entail additional I/O when a frame is chosen for replacement.
(*Write-Ahead Log* protocol; more later.)

When a Page is Requested ...

- ❖ If requested page is not in pool:
 - Choose a frame for *replacement*
 - If frame is dirty, write it to disk
 - Read requested page into chosen frame
- ❖ *Pin* the page and return its address.
- ➡ *If requests can be predicted (e.g., sequential scans) pages can be pre-fetched several pages at a time!*

Buffer Replacement Policy

- ❖ Frame is chosen for replacement by a *replacement policy*:
 - Least-recently-used (LRU), Clock, MRU etc.
- ❖ Policy can have big impact on # of I/O's; depends on the *access pattern*.
- ❖ Sequential flooding: Nasty situation caused by LRU + repeated sequential scans.
 - # buffer frames < # pages in file means each page request causes an I/O. MRU much better in this situation (but not in all situations, of course).

DBMS vs. OS File System

OS does disk space & buffer mgmt: why not let OS manage these tasks?

- ❖ Differences in OS support: portability issues
- ❖ Some limitations, e.g., files can't span disks.
- ❖ Buffer management in DBMS requires ability to:
 - *pin a page* in buffer pool, *force a page* to disk (important for implementing CC & recovery),
 - adjust *replacement policy*, and *pre-fetch pages* based on access patterns in typical DB operations.