

# *Files*

- ❖ Access method layer offers an abstraction of data on disk: **a file of records residing on multiple pages**
  - A number of fields are organized in a record
  - A collection of records are organized in a page
  - A collection of pages are organized in a file

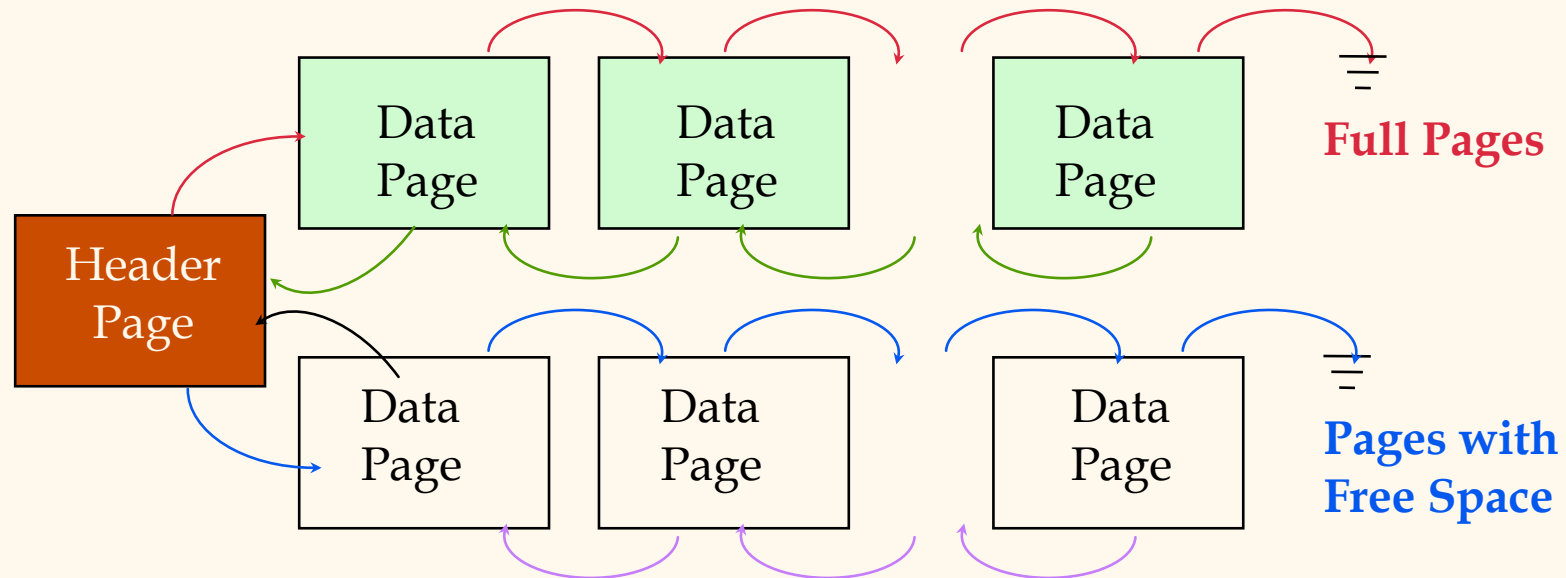
# *Files of Records*

- ❖ Page or block is OK when doing I/O, but higher levels of DBMS operate on *records* and *files of records*.
- ❖ FILE: A collection of pages, each containing a collection of records. Must support:
  - insert/delete/modify record
  - read a particular record (specified using *record id*)
  - scan all records (possibly with some conditions on the records to be retrieved)

# *Unordered (Heap) Files*


- ❖ Simplest file structure contains records in no particular order.
- ❖ As file grows and shrinks, disk pages are allocated and de-allocated.
- ❖ To support record level operations, we must:
  - keep track of the pages in a file
  - keep track of free space on pages
  - keep track of the records on a page
- ❖ There are many alternatives for keeping track of this.

# *Heap File Implemented as a List*



- ❖ (heap file name, header page id) stored in a known place.
- ❖ Two doubly linked lists, for full pages & pages with space.
  - Each page contains 2 'pointers' plus data.
- ❖ Upon insertion, scan the list of pages with space, or ask disk space manager to allocate a new page

# Heap File Using a Page Directory




- ❖ A directory entry per page; it can include the number of free bytes on the page.
- ❖ The directory is a collection of pages; linked list implementation is just one alternative.
  - *Much smaller than linked list of all HF pages!*

# *Page Format*


- ❖ How to store a collection of records on a page?
- ❖ Consider a page as a collection of **slots**, one for each record.
- ❖ A record is identified by **rid = <page id, slot #>**
- ❖ Record ids (rids) are used in indexes (Alternatives 2 and 3).

# Page Format: Fixed Length Records



➡ *Moving records for free space management changes rid! May not be acceptable.*


# Page Format: Variable Length Records



- ➡ Can move records on page without changing rid; so, attractive for fixed-length records too. (*level of indirection*)




# Record Format: Fixed Length



- ❖ Information of a *record type* e.g., the number of fields and field types is stored in the *system catalog*.
- ❖ **Fixed length record:** (1) the number of fields is fixed, (2) each field has a fixed length.
- ❖ Store fields consecutively in a record.
- ❖ Finding *i'th* field does not require scan of record.

# Record Format: Variable Length

- ❖ **Variable length record:** (1) number of fields is fixed, (2) some fields are variable length
- ❖ Two alternatives:



➡ Second offers direct access to  $i$ 'th field, efficient storage of NULLs ; small directory overhead.

# *System Catalogs*

- ❖ For each index:
  - structure (e.g., B+ tree) and search key fields
- ❖ For each relation:
  - name, file name, file structure (e.g., Heap file)
  - attribute name and type, for each attribute
  - index name, for each index
  - integrity constraints
- ❖ For each view:
  - view name and definition
- ❖ Plus statistics, authorization, buffer pool size, etc.

➡ *Catalogs are themselves stored as relations!*