

Tree-Structured Indexes

CMPSCI 445


Fall 2008

Review

- ❖ *As for any index, 3 alternatives for data entries k^* :*
 - Data record with key value k
 - $\langle k, \text{rid of data record with search key value } k \rangle$
 - $\langle k, \text{list of rids of data records with search key } k \rangle$
- ❖ Choice is orthogonal to the *indexing technique* used to locate data entries k^* .
- ❖ Tree-structured indexing techniques support both *range searches* and *equality searches*.


B+ Tree: Most Widely Used Index

- ❖ Inserts/deletes keep tree *height-balanced*. $\log_F N$ cost (F = fanout, N = # leaf pages).
- ❖ **Minimum 50% occupancy** (except for root). Each node contains $d \leq m \leq 2d$ entries, where d is called the *order* of the tree.
- ❖ Supports equality, range-searches, updates efficiently.



Example B+ Tree

- ❖ Search begins at root, and key comparisons direct it to a leaf.
- ❖ Search for 5*, 15*, all data entries $\geq 24^*$...



➡ Based on the search for 15*, we know it is not in the tree!

B+ Trees in Practice


- ❖ Typical **order**: 100. Typical **fill-factor**: 67%.
 - average fanout = 133
- ❖ Typical capacities:
 - Height 4: $133^4 = 312,900,700$ records
 - Height 3: $133^3 = 2,352,637$ records
- ❖ Can often hold top levels in buffer pool:
 - Level 1 = 1 page = 8 Kbytes
 - Level 2 = 133 pages = 1 Mbyte
 - Level 3 = 17,689 pages = 133 MBytes

Inserting a Data Entry into a B+ Tree

- ❖ Find correct leaf L .
- ❖ Put data entry onto L .
 - If L has enough space, *done!*
 - Else, must split L (into L and a new node $L2$)
 - Redistribute entries evenly, copy up middle key.
 - Insert index entry pointing to $L2$ into parent of L .
- ❖ This can happen recursively
 - To split index node, redistribute entries evenly, but push up middle key. (Contrast with leaf splits.)
- ❖ Splits “grow” tree; root split increases height.
 - Tree growth: gets wider or one level taller at top.



Previous Example

Inserting 8*




Inserting 8* into Example B+ Tree

- ❖ Observe how minimum occupancy is guaranteed in both leaf and index pg splits.
- ❖ Note difference between *copy-up* and *push-up*; be sure you understand the reasons for this.



Example B+ Tree After Inserting 8*



- ❖ Notice that root was split, leading to increase in height.
- ❖ In this example, we can avoid split by re-distributing entries between siblings; but not usually done in practice.