

Outline

- ❖ Sorting
- ❖ Evaluation of joins
- ❖ Evaluation of other operations

Some Common Techniques

- ❖ Algorithms for evaluating relational operators use some simple ideas extensively:
 - **Indexing:** Can use WHERE conditions to retrieve small set of tuples (selections, joins)
 - **Iteration:** Sometimes, faster to scan all tuples even if there is an index. (And sometimes, we can scan the data entries in an index instead of the table itself.)
 - **Partitioning:** By using sorting or hashing, we can partition the input tuples and replace an expensive operation by similar operations on smaller inputs.

** Watch for these techniques as we discuss query evaluation!*

Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Reserves (*sid*: integer, *bid*: integer, *day*: date, *rname*: string)

❖ Reserves:

- Each tuple is 40 bytes long,
- 100 tuples per page,
- 1000 pages.

❖ Sailors:

- Each tuple is 50 bytes long,
- 80 tuples per page,
- 500 pages.

Equality Joins With One Join Column

```
SELECT *  
FROM   Reserves R1, Sailors S1  
WHERE  R1.sid=S1.sid
```

- ❖ In algebra: $R \bowtie S$. Common relational operation!
 - $R \times S$ is large; $R \times S$ followed by a selection is inefficient.
 - Must be carefully optimized.
- ❖ Assume: M pages in R , p_R tuples per page, N pages in S , p_S tuples per page.
 - In our examples, R is Reserves and S is Sailors.
- ❖ We will consider more complex join conditions later.
- ❖ *Cost metric*: # of I/Os. We will ignore output costs.

Simple Nested Loops Join

```
foreach tuple r in R do
  foreach tuple s in S do
    if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```


- ❖ For each tuple in the *outer* relation R, we scan the entire *inner* relation S.
 - Cost: $M + p_R * M * N = 1000 + 100 * 1000 * 500 = 1,000 + (5 * 10^7)$ I/Os.
 - Assuming each I/O takes 10 ms, the join will take about 140 hours!

Page-Oriented Nested Loops Join

- ❖ For each *page* of R, get each *page* of S, and write out matching pairs of tuples $\langle r, s \rangle$, where r is in R-page and S is in S-page.
 - **Cost:** $M + M * N = 1000 + 1000*500 = 501,000$ I/Os.
 - Assuming each I/O takes 10 ms, the join will take about 1.4 hours.
- ❖ Choice of the *smaller* relation as the *outer*
 - If smaller relation (S) is outer, cost = $500 + 500*1000 = 500,500$ I/Os.

Block Nested Loops Join

- ❖ Take the smaller relation, say R, as outer, the other as inner.
- ❖ Use one buffer for scanning the inner S, one buffer for output, and use all remaining buffers to hold ``block'' of outer R.
 - For each matching tuple r in R-block, s in S-page, add $\langle r, s \rangle$ to result.
 - Then read next page in S, until S is finished.
 - Then read next R-block, scan S...



Examples of Block Nested Loops

- ❖ **Cost: Scan of outer + #outer blocks * scan of inner**
 - #outer blocks = $\lceil \# \text{ pages of outer} / \text{block size} \rceil$
 - Given available buffer size B, block size is at most B-2.
 - $M + N * \lceil M / B-2 \rceil$
- ❖ With Sailors (S) as outer, a block has 100 pages of S:
 - Cost of scanning S is 500 I/Os; a total of 5 *blocks*.
 - Per block of S, we scan Reserves; 5*1000 I/Os.
 - Total = $500 + 5 * 1000 = 5,500$ I/Os.

Index Nested Loops Join

```
foreach tuple r in R do
    foreach tuple s in S where  $r_i == s_j$  do
        add <r, s> to result
```

- ❖ If there is an index on the join column of one relation (say S), can make it the inner and exploit the index.
 - Cost: $M + (M * p_R) * \text{cost of finding matching S tuples}$
- ❖ For each R tuple, cost of probing S index is about 1.2 for hash index, 2-4 for B+ tree. Cost of then finding S tuples (assuming Alt. (2) or (3) for data entries) depends on clustering.
 - Clustered index: 1 I/O (typical).
 - Unclustered: upto 1 I/O per matching S tuple.

Examples of Index Nested Loops

- ❖ Hash-index (Alt. 2) on *sid* of Sailors (as inner):
 - Scan Reserves: 1000 page I/Os, 100*1000 tuples.
 - For each Reserves tuple: 1.2 I/Os to get data entry in index, plus 1 I/O to get (the exactly one) matching Sailors tuple.
 - Total: $1000 + 100 \cdot 1000 \cdot 2.2 = 221,000$ I/Os.
- ❖ Hash-index (Alt. 2) on *sid* of Reserves (as inner):
 - Scan Sailors: 500 page I/Os, 80*500 tuples.
 - For each Sailors tuple: 1.2 I/Os to find index page with data entries, plus cost of retrieving matching Reserves tuples. If uniform distribution, 2.5 reservations per sailor (100,000 / 40,000). Cost of retrieving them is 1 or 2.5 I/Os (cluster?).
 - Total: $500 + 80 \cdot 500 \cdot (2.2 \sim 3.7) = 88,500 \sim 148,500$ I/Os.

Sort-Merge Join ($R \bowtie_{i=j} S$)

- ❖ (1) Sort R and S on the join column, (2) Merge them (on join col.), and output result tuples.
- ❖ Merge: repeat until either R or S is finished
 - *Scanning*: Advance scan of R until current R-tuple \geq current S tuple, advance scan of S until current S-tuple \geq current R tuple; do this until **current R tuple = current S tuple**.
 - *Matching*: Now all R tuples with same value in R_i (*current R group*) and all S tuples with same value in S_j (*current S group*) match; output $\langle r, s \rangle$ for all pairs of such tuples.
- ❖ R is scanned once; each S group is scanned once per matching R tuple. (Multiple scans of an S group are likely to find needed pages in buffer.)

Example of Sort-Merge Join

<u>sid</u>	sname	rating	age	<u>sid</u>	<u>bid</u>	<u>day</u>	rname
22	dustin	7	45.0	28	103	12/4/96	guppy
28	yuppy	9	35.0	28	103	11/3/96	yuppy
31	lubber	8	55.5	31	101	10/10/96	dustin
44	guppy	5	35.0	31	102	10/12/96	lubber
58	rusty	10	35.0	31	101	10/11/96	lubber
				58	103	11/12/96	dustin

❖ Cost: $M \log M + N \log N + (M+N)$

- The cost of merging, $M+N$, could be $M*N$ (very unlikely!)
- $M+N$ is guaranteed in *foreign key join* (why?)
- As with sorting, $\log M$ and $\log N$ are small numbers, e.g., 3, 4.

❖ With 35, 100 or 300 buffer pages, both Reserves and Sailors can be sorted in 2 passes; total join cost: 7500.

(BNL cost: 2500 ($B=300$), 5500 ($B=100$), 15000 ($B=35$))