

Scheduling Transactions

- ❖ Serial schedule: Schedule that does not interleave the actions of different transactions.
- ❖ Equivalent schedules: For any database state, the effect (on the set of objects in the database) of executing the first schedule is identical to the effect of executing the second schedule.
- ❖ Serializable schedule: A schedule that is equivalent to some serial execution of the transactions.

Serializable Schedule

T1	T2
R(A) W(A)	
	R(A) W(A)
R(B) W(B)	
	R(B) W(B) Commit
Commit	

When can actions be re-ordered?

- ❖ Let I, J be two consecutive actions of T1 and T2
 - I=Read(O), J=Read(O)
 - I=Read(O), J=Write(O)
 - I=Write(O), J=Read(O)
 - I=Write(O), J=Write(O)
- ❖ If I and J are both reads, then they can be freely reordered.
- ❖ In all other cases, order impacts outcome of schedule.

Conflicting operations

- ❖ Two operations **conflict** if:
 - they operate on the same data object, and
 - at least one is a WRITE.
- ❖ Schedule outcome is determined by order of the conflicting operations.

Conflict Serializable Schedules

- ❖ Two schedules are **conflict equivalent** if:
 - Involve the same actions of the same transactions
 - Every pair of conflicting actions (of committed trans) are ordered the same way.
 - Alternatively: S can be transformed to S' by swaps of non-conflicting actions.
- ❖ Schedule S is **conflict serializable** if S is conflict equivalent to some serial schedule

Every conflict serializable schedule is serializable.

(exception: dynamic databases)

Conflict-serializable schedule

T1	T2
R(A)	
W(A)	
	R(A)
	W(A)
R(B)	
W(B)	
	R(B)
	W(B)
	Commit
Commit	

Not conflict-serializable

T1	T2
R(A) W(A)	R(A) W(A) R(B) W(B) Commit
R(B) W(B) Commit	

Precedence graphs

- ❖ Directed graph derived from schedule S:
 - Vertex for each transaction
 - Edge from T_i to T_j if:
 - T_i executes Write(O) before T_j executes Read(O)
 - T_i executes Read(O) before T_j executes Write(O)
 - T_i executes Write(O) before T_j executes Write(O)

If edge $T_i \rightarrow T_j$ appears in precedence graph, then in any serial schedule equivalent to S, T_i must appear before T_j .

Dependency Graph

- ❖ Theorem: A schedule is **conflict serializable** if and only if its dependency graph is acyclic.

(A serializable order can be found by topological sort of the dependency graph.)


Construct precedence graphs:

T1	T2
R(A) W(A)	
	R(A) W(A)
R(B) W(B)	
	R(B) W(B) Commit
Commit	



Conflict serializable

T1	T2
R(A) W(A)	
	R(A) W(A) R(B) W(B) Commit
R(B) W(B) Commit	



Non-conflict serializable

Construct precedence graph:

T1	T2	T3
R(A)	W(A) Commit	
W(A) Commit		W(A) Commit

Recoverable schedules

- ❖ We must also consider the impact of transaction failures on concurrently running transactions.
 - That is, schedules with ABORT
- ❖ **Recoverable schedule:** if T_j reads data written by T_i , then T_i commits before T_j commits.

A Non-recoverable schedule.

T1	T2
R(A) W(A)	
	R(A) W(A) R(B) W(B) Commit
Abort	

DBMS must ensure recoverable schedules.

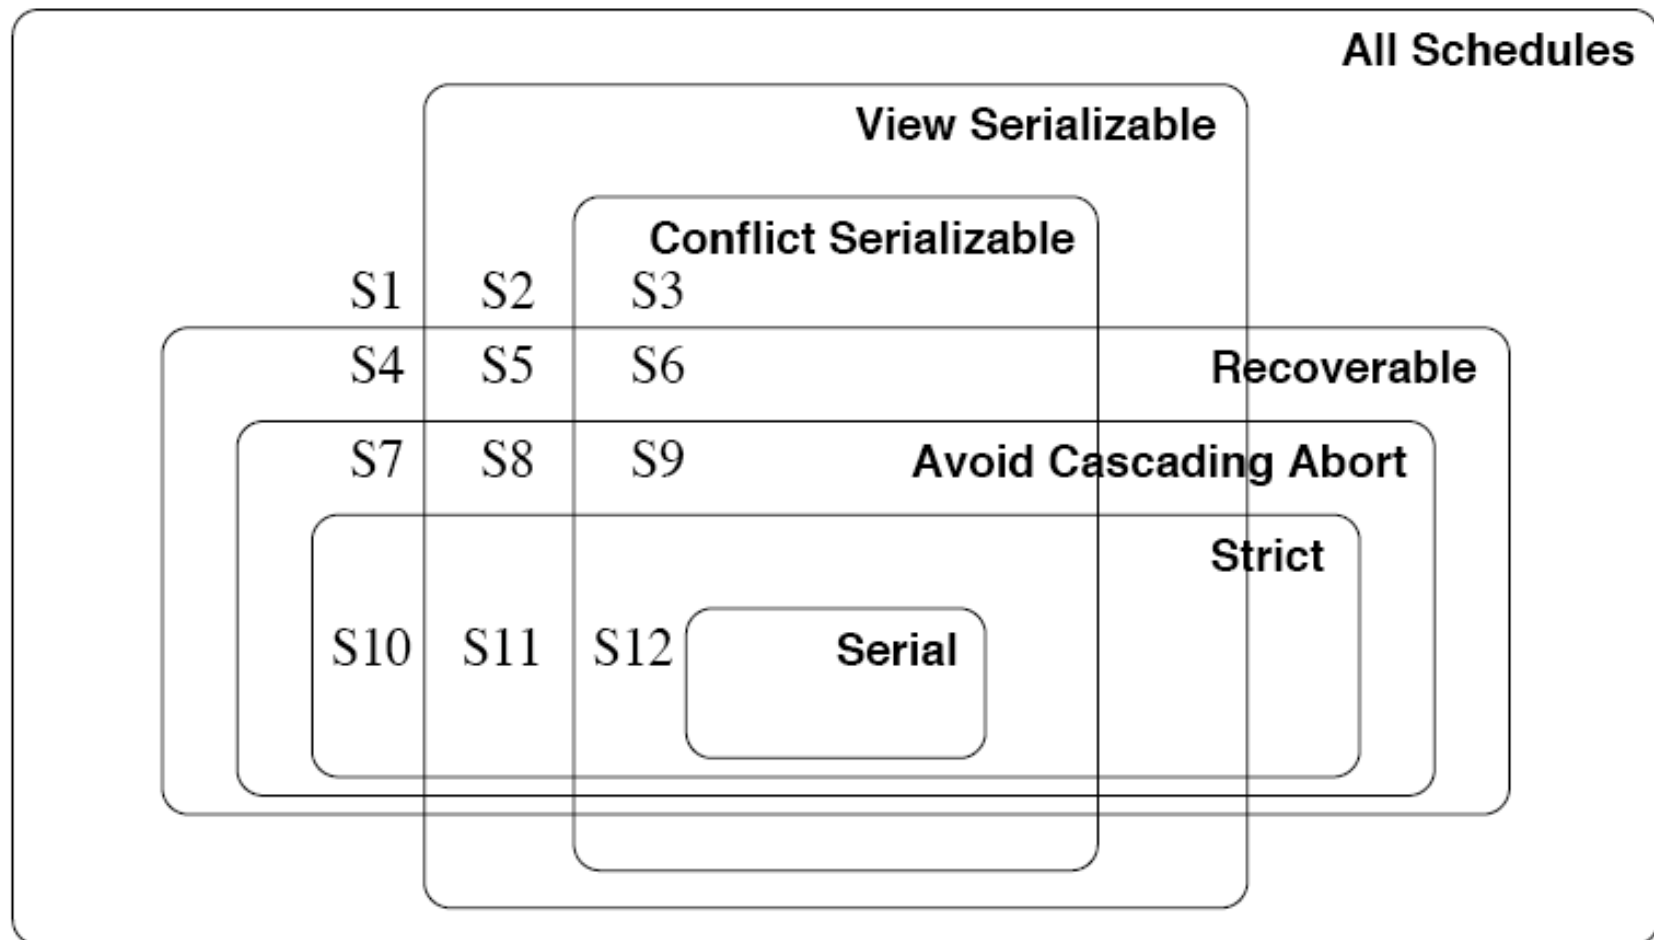
Cascadeless schedules

- ❖ Even if schedule is recoverable, several transactions may need to be rolled back to recover correctly.
- ❖ **Cascading Rollback:** a single transaction failure leading to a series of rollbacks

T1	T2	T3
R(A)		
R(B)		
W(A)		
	R(A)	
	W(A)	
		R(A)
Abort		

- ❖ **Cascadeless schedule:** if T_j reads data written by T_i , then T_i commits before read operation of T_j .

Properties of schedules



Anomalies with Interleaved Execution

- ❖ Not all interleavings of operations are okay.
- ❖ **Anomaly**: two consistency-preserving committed transactions that lead to an inconsistent state.
- ❖ Types of anomalies:
 - Reading Uncommitted Data (WR Conflicts)
“dirty reads”
 - Unrepeatable Reads (RW Conflicts)
 - Overwriting Uncommitted Data (WW Conflicts)

Reading Uncommitted Data

“Dirty Read”

Inconsistent result of A is
exposed to transaction T2

T1: Transfer	T2: Interest
R(A) W(A)	
	R(A) W(A) R(B) W(B) Commit
R(B) W(B) Commit	

Acceptable Dirty Read

If we are just checking availability of an airline seat, a dirty read might be fine! (*Why is that?*)

Reservation transaction:

```
EXEC SQL select occupied into :occ
        from Flights
        where Num= '123' and date=11-03-99
              and seat='23f';
if (!occ) {EXEC SQL
        update Flights
        set occupied=true
        where Num= '123' and date=11-03-99
              and seat='23f';}
else {notify user that seat is unavailable}
```

Unrepeatable Reads (RW Conflicts)

T1 could see two values for A,
although it has not changed A itself.

T1	T2
R(A)	R(A) W(A) Commit
R(A) W(A) Commit	

Overwriting Uncommitted Data

T1	T2
W(B)	W(A)
	W(B)
	Commit
W(A)	
Commit	

Schedules involving abort

T1	T2
R(A) W(A)	R(A) W(A) R(B) W(B) Commit
Abort	

This is an
unrecoverable
schedule

- ❖ Recoverable schedule: transactions commit only after all transactions whose changes they read commit.

Concurrency control schemes

- ❖ The DBMS must provide a mechanism that will ensure all possible schedules are:
 - serializable
 - recoverable, and preferably cascadeless
- ❖ Concurrency control protocols ensure these properties.

Lock-Based Concurrency Control

- ❖ Lock - associated with some object
 - shared or exclusive
- ❖ Locking protocol - set of rules to be followed by each transaction to ensure good properties.

Lock Compatibility Matrix

Locks on a data item are granted based on a
lock compatibility matrix:

		Mode of Data Item		
		None	Shared	Exclusive
Request mode {	Shared	Y	Y	N
	Exclusive	Y	N	N

When a transaction requests a lock, it must wait
(block) until the lock is granted

Strict Two Phase Locking

❖ (Strict 2PL) Protocol:

- Each Xact must obtain a *S (shared)* lock on object before reading,
- Each Xact must obtain an *X (exclusive)* lock on object before writing.
- All locks held by a transaction are released when the transaction completes
- If an Xact holds an X lock on an object, no other Xact can get a lock (S or X) on that object.

❖ Strict 2PL allows only serializable schedules.

- Allows only safe interleavings
- No anomalies

Schedule following strict 2PL

T1	T2
S(A) R(A)	
	S(A) R(A) X(B) R(B) W(B) Commit
X(C) R(C) W(C) Commit	

Deadlock

T1	T2
X(A)	
	X(B)
X(B)	
	X(A)

granted

granted

queued

queued

❖ Deadlock must be prevented or avoided.

Performance of Locking

- ❖ Lock-based schemes resolve conflicting schedules by **blocking** and **aborting**
 - in practice few deadlocks and relatively few aborts
 - most of penalty from blocking
- ❖ To increase throughput
 - lock smallest objects possible
 - reduce time locks are held
 - reduce hotspots

Transaction support in SQL

- ❖ Transaction automatically started for SELECT, UPDATE, CREATE
- ❖ Transaction ends with COMMIT or ROLLBACK (abort)
- ❖ SQL 99 supports SAVEPOINTS which are simple nested transactions

What should we lock?

T1

```
SELECT S.rating, MIN(S.age)
FROM Sailors S
WHERE S.rating = 8
```

T2

```
UPDATE Sailors(Name, Rating,
Age) VALUES ("Joe", 8, 33)
```

- ❖ T1 **S**-lock on Sailors; T2 **X**-lock on Sailors
- ❖ T1 **S**-lock on all rows with rating=8; T2 **X**-lock on Joe's tuple.

The Phantom Problem

- ❖ T1 locks all **existing** rows with rating=8.
- ❖ But a new row satisfying condition could be inserted
- ❖ Phantom problem: A transaction retrieves a collection of tuples and sees different results, even though it did not modify the tuples itself.
 - Conceptually: must lock all *possible* rows.
 - Can lock entire table.
 - Better, use index locking.

The Phantom Problem

T1

```
SELECT S.rating, MIN(S.age)
FROM Sailors S
WHERE S.rating = 8
```

T2

```
UPDATE Sailors(Name, Rating,
Age) VALUES ("Joe", 8, 33)
```

T3

```
INSERT Sailors(Name, Rating,
Age) VALUES ("Mary", 8, 18)
```

- ❖ Suppose T1 locks all **existing** rows with rating=8.
- ❖ Then T3 creates row and sets X-lock
- ❖ T1 returns an answer that depends on its order relative to T3, but locking does not impose a relative order on these transactions.

Specify isolation level

- General rules of thumb w.r.t. isolation:
 - Fully serializable isolation is more expensive than “no isolation”
 - We can’t do as many things concurrently (or we have to undo them frequently)
- ❖ For performance, we generally want to specify the most relaxed **isolation level** that’s acceptable
 - Note that we’re “slightly” violating a correctness constraint to get performance!

Specifying isolation level in SQL

SET TRANSACTION [READ WRITE | READ ONLY]
ISOLATION LEVEL [LEVEL];

LEVEL =	SERIALIZABLE	↓ Less isolation
	REPEATABLE READ	
	READ COMMITTED	
	READ UNCOMMITTED	

The default isolation level is **SERIALIZABLE**

Locks sets of objects, avoids phantoms

REPEATABLE READ

- ❖ T reads only changes made by committed transactions
- ❖ No value read/written by T is changed by another transaction until T completes.
- ❖ Phantoms possible: inserts of qualifying tuples not avoided.

Locks only individual objects

READ COMMITTED

- ❖ T reads only changes made by committed transactions
- ❖ No value ~~read~~/written by T is changed by another transaction until T completes.
- ❖ Value read by T may be modified while T in progress.
- ❖ Phantoms possible.

X locks on written objects, held to end
S locks on read objects, but released immediately.

READ UNCOMMITTED

- ❖ Greatest exposure to other transactions
- ❖ Dirty reads possible
- ❖ Can't make changes: must be READ ONLY
- ❖ Does not obtain shared locks before reading
 - Thus no locks every requested.

Specifying Acceptable Isolation Levels

- ❖ To signal to the system that a **dirty read** is acceptable,

```
SET TRANSACTION READ WRITE  
ISOLATION LEVEL READ UNCOMMITTED;
```

- ❖ In addition, there are

- ❖

```
SET TRANSACTION  
ISOLATION LEVEL READ COMMITTED;  
SET TRANSACTION  
ISOLATION LEVEL REPEATABLE READ;
```

Summary of Isolation Levels

Level	Dirty Read	Unrepeatable Read	Phantoms
READ UN-COMMITTED	Maybe	Maybe	Maybe
READ COMMITTED	No	Maybe	Maybe
REPEATABLE READ	No	No	Maybe
SERIALIZABLE	No	No	No

Summary

- ❖ Concurrency control and recovery are among the most important functions provided by a DBMS.
- ❖ Users need not worry about concurrency.
 - System guarantees nice properties: ACID
 - This is implemented using a locking protocol
- Users can trade isolation for performance using SQL commands