

# *Final concurrency topics, Crash Recovery*

Fall 2008

Slide content adapted from Ramakrishnan & Gehrke

# *Lock-Based Concurrency Control*

- ❖ DBMS must ensure
  - only serializable, recoverable schedules are allowed
  - No actions of committed trans lost while undoing aborted trans
- ❖ Lock - associated with some object
  - shared or exclusive
- ❖ Locking protocol - set of rules to be followed by each transaction to ensure good properties.

# *Two-Phase Locking (2PL)*

## ❖ Two-Phase Locking Protocol

- Each Xact must obtain a S (*shared*) lock on object before reading, and an X (*exclusive*) lock on object before writing.
- A transaction can not request additional locks once it releases any locks.
- If an Xact holds an X lock on an object, no other Xact can get a lock (S or X) on that object.

## ❖ 2PL ensures **conflict serializability**

- Transactions can be ordered by their end of growing phase; this is serializable order.

# Strict Two Phase Locking (Strict 2PL)

## ❖ Strict Two-phase Locking (Strict 2PL) Protocol:

- Each Xact must obtain a *S (shared)* lock on object before reading, and an *X (exclusive)* lock on object before writing.
- All locks held by a transaction are released when the transaction completes.
- If an Xact holds an X lock on an object, no other Xact can get a lock (S or X) on that object
- Strict 2PL ensures **conflict serializable** and **cascadeless** schedules

# *What should we lock?*

**T1**

```
SELECT S.rating, MIN(S.age)
FROM Sailors S
WHERE S.rating = 8
```

**T2**

```
UPDATE Sailors(Name, Rating,
Age) VALUES ("Joe", 8, 33)
```

- ❖ T1 **S**-lock on Sailors; T2 **X**-lock on Sailors
- ❖ T1 **S**-lock on all rows with rating=8; T2 **X**-lock on Joe's tuple.

# *Phantom*

- ❖ T1: “Find oldest sailor for each of the rating levels 1 and 2”
  - T1 locks all pages containing sailor records with *rating* = 1, and finds oldest sailor (say, *age* = 71).
- ❖ T2: “Insert new sailor. *rating*=1, *age*=96”
- ❖ T2: “Deletes oldest sailor with *rating* = 2 (and, say, *age* = 80), and commits
- ❖ T1 now locks all pages containing sailor records with *rating* = 2, and finds oldest (say, *age* = 63).

# *The Problem*

- ❖ T1 implicitly assumes that it has locked the set of all sailor records with *rating* = 1.
  - Assumption only holds if no sailor records are added while T1 is executing!
  - Need some mechanism to enforce this assumption. (Index locking and predicate locking.)
- ❖ Example shows that conflict serializability guarantees serializability only if the set of objects is fixed!
- ❖ Strict 2PL will not assure serializability

# *The Phantom Problem*

- ❖ Phantom problem: A transaction retrieves a collection of tuples and sees different results, even though it did not modify the tuples itself.
  - Conceptually: must lock all *possible* rows.
  - Can lock entire table.
  - Better, use index locking.



# *Specify isolation level*

- General rules of thumb w.r.t. isolation:
  - Fully serializable isolation is more expensive than “no isolation”
    - We can’t do as many things concurrently (or we have to undo them frequently)
- ❖ For performance, we generally want to specify the most relaxed **isolation level** that’s acceptable for our application.
  - Note that we’re “slightly” violating a correctness constraint to get performance!

# *Specifying isolation level in SQL*

SET TRANSACTION [READ WRITE | READ ONLY]  
ISOLATION LEVEL [LEVEL];

LEVEL =	SERIALIZABLE	↓ Less isolation
	REPEATABLE READ	
	READ COMMITTED	
	READ UNCOMMITTED	

The default isolation level is **SERIALIZABLE**

Locks sets of objects, avoids phantoms

# *REPEATABLE READ*

- ❖ T reads only changes made by committed transactions
- ❖ No value read/written by T is changed by another transaction until T completes.
- ❖ Phantoms possible: inserts of qualifying tuples not avoided.

Locks only individual objects

# *READ COMMITTED*

- ❖ T reads only changes made by committed transactions
- ❖ No value ~~read~~/written by T is changed by another transaction until T completes.
- ❖ Value read by T may be modified while T in progress.
- ❖ Phantoms possible.

X locks on written objects, held to end

S locks on read objects, but released immediately.

# *READ UNCOMMITTED*

- ❖ Greatest exposure to other transactions
- ❖ Dirty reads possible
- ❖ Can't make changes: must be READ ONLY
- ❖ Does not obtain shared locks before reading
  - Thus no locks ever requested.

# *Summary of Isolation Levels*

Level	Dirty Read	Unrepeatable Read	Phantoms
READ UN-COMMITTED	Maybe	Maybe	Maybe
READ COMMITTED	No	Maybe	Maybe
REPEATABLE READ	No	No	Maybe
SERIALIZABLE	No	No	No

# *Recovery*

## *Review: The ACID properties*

- ❖ **A**tomicity: All actions in the Xact happen, or none happen.
- ❖ **C**onsistency: If each Xact is consistent, and the DB starts consistent, it ends up consistent.
- ❖ **I**solation: Execution of one Xact is isolated from that of other Xacts.
- ❖ **D**urability: If a Xact commits, its effects persist.
- ❖ The **Recovery Manager** guarantees **Atomicity & Durability**.



# *Types of failure*

## ❖ Transaction failure

- partially-executed transaction cannot commit
- ➔ changes must be removed: ROLLBACK


## ❖ System failure

- volatile memory lost
- ➔ updates of committed Xact persist
- ➔ updates of aborted or partial Xacts removed

## ❖ Media failure

- corrupted storage media
- ➔ database brought up-to-date using backup

# Motivation



- ❖ Desired Behavior after system restarts:
  - T1, T2 & T3 should be **durable**.
  - T4 & T5 should be **aborted** (effects not seen).

# *Undo and Redo*

## ❖ UNDO:

- removing effects of incomplete or aborted transaction (for atomicity)

## ❖ REDO:

- re-instating effects of committed transactions (for durability)

- ❖ The work the recovery subsystem must do to support UNDO and REDO depends on **key policies** of the buffer manager.

# *Handling the Buffer Pool*


- ❖ **Force** every write to disk?
  - Poor response time.
  - But provides durability.
- ❖ **Steal** buffer-pool frames from uncommitted Xacts?
  - If not, poor throughput.
  - If so, how can we ensure atomicity?

	No Steal	Steal
Force	Trivial	
No Force		Desired

## *More on Steal and Force*

- ❖ STEAL (why enforcing Atomicity is hard)
  - *To steal frame F*: Current page in F (say P) is written to disk; some Xact holds lock on P.
    - What if the Xact with the lock on P aborts?
    - Must remember the old value of P at steal time (to support **UNDO**ing the write to page P).
- ❖ NO FORCE (why enforcing Durability is hard)
  - What if system crashes before a modified page is written to disk?
  - Write as little as possible, in a convenient place, at commit time, to support **REDO**ing modifications.

# Basic Idea: Logging



- ❖ Record REDO and UNDO information, for every update, in a *log*.
  - Sequential writes to log (put it on a separate disk).
  - Minimal info (diff) written to log, so multiple updates fit in a single log page.
- ❖ Log: An ordered list of REDO/UNDO actions
  - Log record contains:
    - Before image (for UNDO), After image (for REDO)
  - and additional control info (which we'll see soon).

# Write-Ahead Logging (WAL)

- ❖ The Write-Ahead Logging Protocol:

- ① Must force the **log record** for an update before the corresponding **data page** is overwritten on disk.
- ② Must **write all log records** for a Xact before commit.

- ❖ #1 guarantees Atomicity.

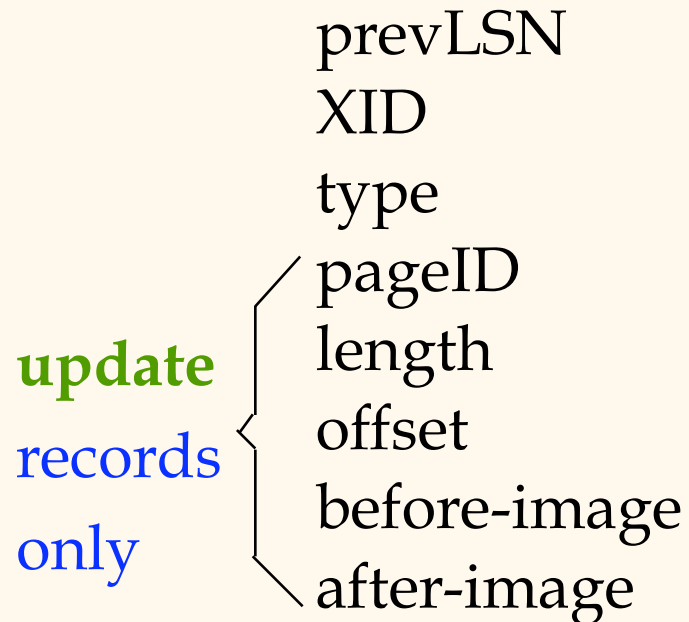
- ❖ #2 guarantees Durability.

- ❖ Exactly how is logging and recovery done?

- We'll study the ARIES algorithms.

# Log Records

## LogRecord fields:



Possible log record types:

- ❖ Update
- ❖ Commit
- ❖ Abort
- ❖ End (signifies end of commit or abort)
- ❖ Compensation Log Records (CLRs)
  - for UNDO actions



# *Other Log-Related State*

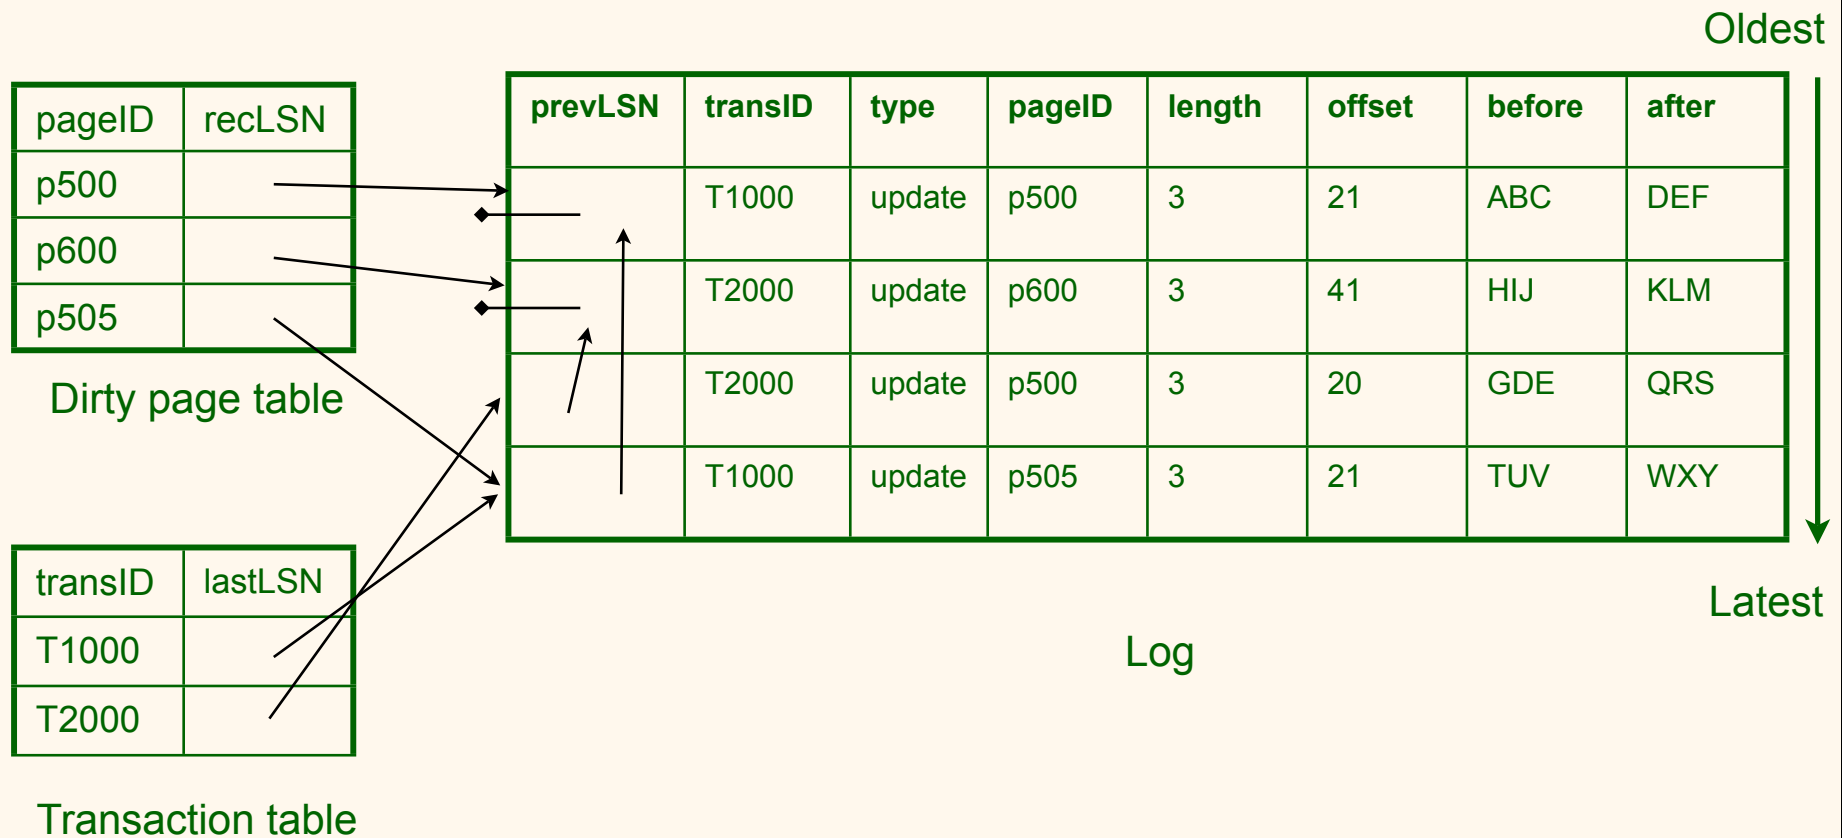
## ❖ Transaction Table:

- One entry per active Xact.
- Contains **XID**, **status** (running/committed/aborted), and **lastLSN**.

## ❖ Dirty Page Table:

- One entry per dirty page in buffer pool.
- Contains **recLSN** -- the LSN of the log record which first caused the page to be dirty.

# Log and Transaction table



# Checkpointing

- ❖ Periodically, the DBMS creates a checkpoint, in order to minimize the time taken to recover in the event of a system crash. Write to log:
  - **begin\_checkpoint** record: Indicates when chkpt began.
  - **end\_checkpoint** record: Contains current *Xact table* and *dirty page table*. This is a '**fuzzy checkpoint**':
    - Other Xacts continue to run; so these tables accurate only as of the time of the **begin\_checkpoint** record.
    - No attempt to force dirty pages to disk; effectiveness of checkpoint limited by oldest unwritten change to a dirty page. (So it's a good idea to periodically flush dirty pages to disk!)
  - Store LSN of chkpt record in a safe place (*master* record).

# *The Big Picture: What's Stored Where*



## **LogRecords**

prevLSN  
XID  
type  
pageID  
length  
offset  
before-image  
after-image



**Data pages**  
each  
with a  
pageLSN  
**master record**



## **Xact Table**


lastLSN  
status

## **Dirty Page Table**

recLSN

**flushedLSN**

# Crash Recovery: Big Picture



- ❖ Start from a **checkpoint** (found via **master** record).
- ❖ Three phases. Need to:
  - **Analysis**: Figure out which Xacts committed since checkpoint, which failed.
  - **REDO** *all* actions.
    - ◆ (repeat history)
  - **UNDO** effects of failed Xacts.

## *Recovery: The Analysis Phase*

- ❖ Reconstruct state of most recent checkpoint.
- ❖ Scan log forward from checkpoint.
  - **End** record: Remove Xact from Xact table.
  - **Other records**: Add Xact to Xact table, set lastLSN=LSN, change Xact status on **commit**.
  - **Update** record: If P not in Dirty Page Table,
    - Add P to D.P.T., set its **recLSN=LSN**.

# Recovery: The REDO Phase

- ❖ We *repeat History* to reconstruct state at crash:
  - Reapply *all* updates (even of aborted Xacts!), redo CLR's.
- ❖ Scan forward from log rec containing smallest *recLSN* in D.P.T. For each CLR or update log rec *LSN*, REDO the action unless:
  1. Affected page is not in the Dirty Page Table, or
  2. Affected page is in D.P.T., but has  $\text{recLSN} > \text{LSN}$ , or
  3.  $\text{pageLSN (in DB)} \geq \text{LSN}$ .
- ❖ To *REDO* an action:
  - Reapply logged action.
  - Set *pageLSN* to *LSN*. No additional logging!

# *Recovery: The UNDO Phase*

ToUndo = {  $l$  |  $l$  a lastLSN of a “loser” Xact }

**Repeat:**

- Choose largest LSN among ToUndo.
- If this LSN is a CLR and undonextLSN == NULL
  - Write an End record for this Xact.
- If this LSN is a CLR, and undonextLSN != NULL
  - Add undonextLSN to ToUndo
- Else this LSN is an update. Undo the update, write a CLR, add prevLSN to ToUndo.

**Until ToUndo is empty.**



# *Summary of Logging/Recovery*

- ❖ **Recovery Manager** guarantees Atomicity & Durability.
- ❖ Use WAL to allow STEAL/NO-FORCE w/o sacrificing correctness.
- ❖ LSNs identify log records; linked into backwards chains per transaction (via prevLSN).
- ❖ pageLSN allows comparison of data page and log records.

## *Summary, Cont.*

- ❖ **Checkpointing:** A quick way to limit the amount of log to scan on recovery.
- ❖ Recovery works in 3 phases:
  - **Analysis:** Forward from checkpoint.
  - **Redo:** Forward from oldest recLSN.
  - **Undo:** Backward from end to first LSN of oldest Xact alive at crash.
- ❖ Upon Undo, write CLR's.
- ❖ Redo “repeats history”: Simplifies the logic!